

Yandex

RDDs

Why do we need a new abstraction?

Why do we need a new abstraction?

- › **Example:** iterative computations (K-means, PageRank, ...)
 - › relation between consequent steps is known only to the user code
 - › framework must reliably persist data between steps (even if it is temporary data)

Why do we need a new abstraction?

- › **Example:** iterative computations (K-means, PageRank, ...)
 - › relation between consequent steps is known only to the user code
 - › framework must reliably persist data between steps (even if it is temporary data)
- › **Example:** joins
 - › join operation is used in many MapReduce applications
 - › not-so-easy to reuse code

Resilient Distributed Datasets

Resilient Distributed Datasets

- › **Resilient** — able to withstand failures
- › **Distributed** — spanning across multiple machines
- › Formally, a read-only, partitioned collection of records

Resilient Distributed Datasets

- › **Resilient** — able to withstand failures
- › **Distributed** — spanning across multiple machines
- › Formally, a read-only, partitioned collection of records
- › To adhere to RDD interface, a dataset must implement:

Resilient Distributed Datasets

- › **Resilient** — able to withstand failures
- › **Distributed** — spanning across multiple machines
- › Formally, a read-only, partitioned collection of records
- › To adhere to RDD interface, a dataset must implement:
 - › `partitions() → Array[Partition]`
 - › `iterator(p: Partition) → Iterator`

Resilient Distributed Datasets

- › **Resilient** — able to withstand failures
- › **Distributed** — spanning across multiple machines
- › Formally, a read-only, partitioned collection of records
- › To adhere to RDD interface, a dataset must implement:
 - › `partitions() → Array[Partition]`
 - › `iterator(p: Partition, parents: Array[Iterator]) → Iterator`
 - › `dependencies() → Array[Dependency]`

Resilient Distributed Datasets

- › **Resilient** — able to withstand failures
- › **Distributed** — spanning across multiple machines
- › Formally, a read-only, partitioned collection of records
- › To adhere to RDD[**T**] interface, a dataset must implement:
 - › partitions() → Array[Partition]
 - › iterator(p: Partition, parents: Array[Iterator[_]]) → Iterator[**T**]
 - › dependencies() → Array[Dependency]
- › ...and may implement other helper functions
- › Typed! RDD[**T**] — a dataset of items of type **T**

Example: a binary file in HDFS

- › `partitions()` → *Array[Partition]*
- › `iterator(p: Partition, parents: Array[Iterator[_]])` →
→ *Iterator[Byte]*
- › `dependencies()` → *Array[Dependency]*

File



Example: a binary file in HDFS


- › **partitions()** → *Array[Partition]*
 - › lookup blocks information from the NameNode
 - › make a partition for every block
 - › return an array of the partitions
- › **iterator**(*p: Partition*, **parents**: *Array[Iterator[_]]*) →
→ *Iterator[Byte]*
- › **dependencies()** → *Array[Dependency]*

File




Example: a binary file in HDFS

- › **partitions()** → *Array[Partition]*
 - › lookup blocks information from the NameNode
 - › make a partition for every block
 - › return an array of the partitions
- › **iterator**(*p: Partition*, **parents**: *Array[Iterator[_]]*) → *Iterator[Byte]*
 - › parents are not used
- › **dependencies()** → *Array[Dependency]*
 - › return an empty array




Example: a binary file in HDFS

- › **partitions()** → *Array[Partition]*
 - › lookup blocks information from the NameNode
 - › make a partition for every block
 - › return an array of the partitions
- › **iterator**(*p: Partition*, **parents**: *Array[Iterator[_]]*) → *Iterator[Byte]*
 - › parents are not used
 - › return a reader for the block of the given partition
- › **dependencies()** → *Array[Dependency]*
 - › return an empty array



Example: a ~~binary~~ **data*** file in HDFS


- › **partitions()** → *Array[Partition]*
 - › lookup blocks information from the NameNode
 - › make a partition for every block
 - › return an array of the partitions
- › **iterator**(*p: Partition*, **parents**: *Array[Iterator[_]]*) → *Iterator[Byte]*
 - › parents are not used
 - › return a reader for the block of the given partition
- › **dependencies()** → *Array[Dependency]*
 - › return an empty array



*a file encoded with the file format: see W1; think: text file, SequenceFile, Avro, RCFile

Example: a ~~binary~~ **data*** file in HDFS


- › **partitions()** → *Array[Partition]*
 - › ~~lookup blocks information from the NameNode~~
use **InputFormat** to compute **InputSplits**
 - › make a partition for every ~~block~~ **InputSplit**
 - › return an array of the partitions
- › **iterator**(*p: Partition, parents: Array[Iterator[_]]*) → *Iterator[Byte **InputRecord**]*
 - › parents are not used
 - › use **InputFormat** to create a reader for the **InputSplit** of the given partition
 - › ~~return a reader for the block of the given partition~~ **the reader**
- › **dependencies()** → *Array[Dependency]*
 - › return an empty array



*a file encoded with the file format: see W1; think: text file, SequenceFile, Avro, RCFile


Example: an in-memory array

- › `partitions()` \rightarrow *Array[Partition]*
- › `iterator(p: Partition, parents: Array[Iterator[_]])` \rightarrow *Iterator[T]*
- › `dependencies()` \rightarrow *Array[Dependency]*




Example: an in-memory array

- › **partitions()** \rightarrow *Array[Partition]*
 - › return an array of a single partition with the source array
- › **iterator**(*p: Partition*, **parents**: *Array[Iterator[_]]*) \rightarrow *Iterator[T]*
 - › parents are not used
 - › return an iterator over the source array in the given partition
- › **dependencies()** \rightarrow *Array[Dependency]*
 - › return an empty array (no dependencies)



Example: an **sliced*** in-memory array

- › **partitions()** \rightarrow *Array[Partition]*
 - › slice array in chunks of size N
 - › make a partition for every chunk
 - › return an array ~~of a single partition with the source array~~ **of the partitions**
- › **iterator**(*p: Partition*, **parents**: *Array[Iterator[_]]*) \rightarrow *Iterator[T]*
 - › parents are not used
 - › return an iterator over the source array **chunk** in the given partition
- › **dependencies()** \rightarrow *Array[Dependency]*
 - › return an empty array (no dependencies)



*can be used to parallelize in-memory computations

Quiz

Summary

- › RDD is a read-only, partitioned collection of records
 - › a developer can access the partitions and create iterators over them
 - › RDD tracks dependencies (to be explained in the next video)
- › Examples of RDDs
 - › Hadoop files with the proper file format
 - › In-memory arrays
- › Next video: Transformations

BigDATAteam