

**Y**andex


Execution & scheduling

# SparkContext

- › Tells your application how to access a cluster
- › Coordinates processes on the cluster to run your application


# SparkContext

- › Tells your application how to access a cluster
- › Coordinates processes on the cluster to run your application




# SparkContext

- › Tells your application how to access a cluster
- › Coordinates processes on the cluster to run your application



# SparkContext

- › Tells your application how to access a cluster
- › Coordinates processes on the cluster to run your application



# Jobs, stages, tasks

- › **Task** is a unit of work to be done
- › Tasks are created by a **job scheduler** for every job stage
- › **Job** is spawned in response to a Spark action
- › Job is divided in smaller sets of tasks called stages

# Jobs, stages, tasks (example)

›  $Z = X$   
  .map(lambda x: (x % 10, x / 10))  
  .reduceByKey(lambda x, y: x + y)



# Jobs, stages, tasks (example)

```
› Z = X  
  .map(lambda x: (x % 10, x / 10))  
  .reduceByKey(lambda x, y: x + y)  
  .collect()
```

1. Invoking an action...

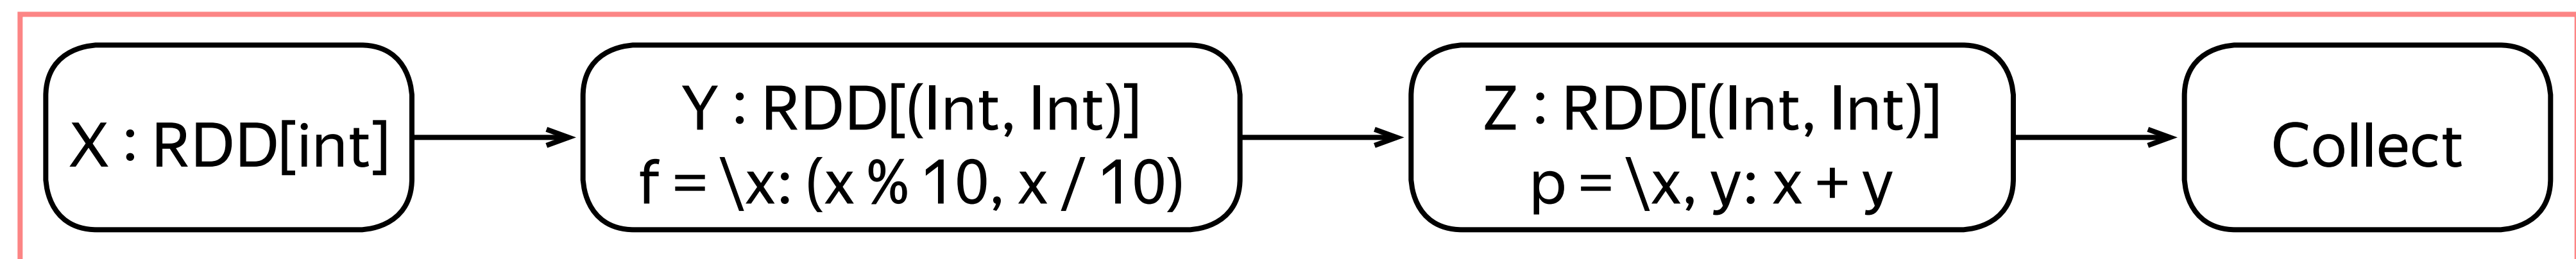
# Jobs, stages, tasks (example)

›  $Z = X$

```
.map(lambda x: (x % 10, x / 10))  
.reduceByKey(lambda x, y: x + y)  
.collect()
```

1. Invoking an action...


2. ...spawns the job...



# Jobs, stages, tasks (example)

›  $Z = X$   
  .map(lambda x: (x % 10, x / 10))  
  .reduceByKey(lambda x, y: x + y)  
  .collect()

1. Invoking an action...
2. ...spawns the job...
3. ...that gets divided into the stages by the job scheduler...

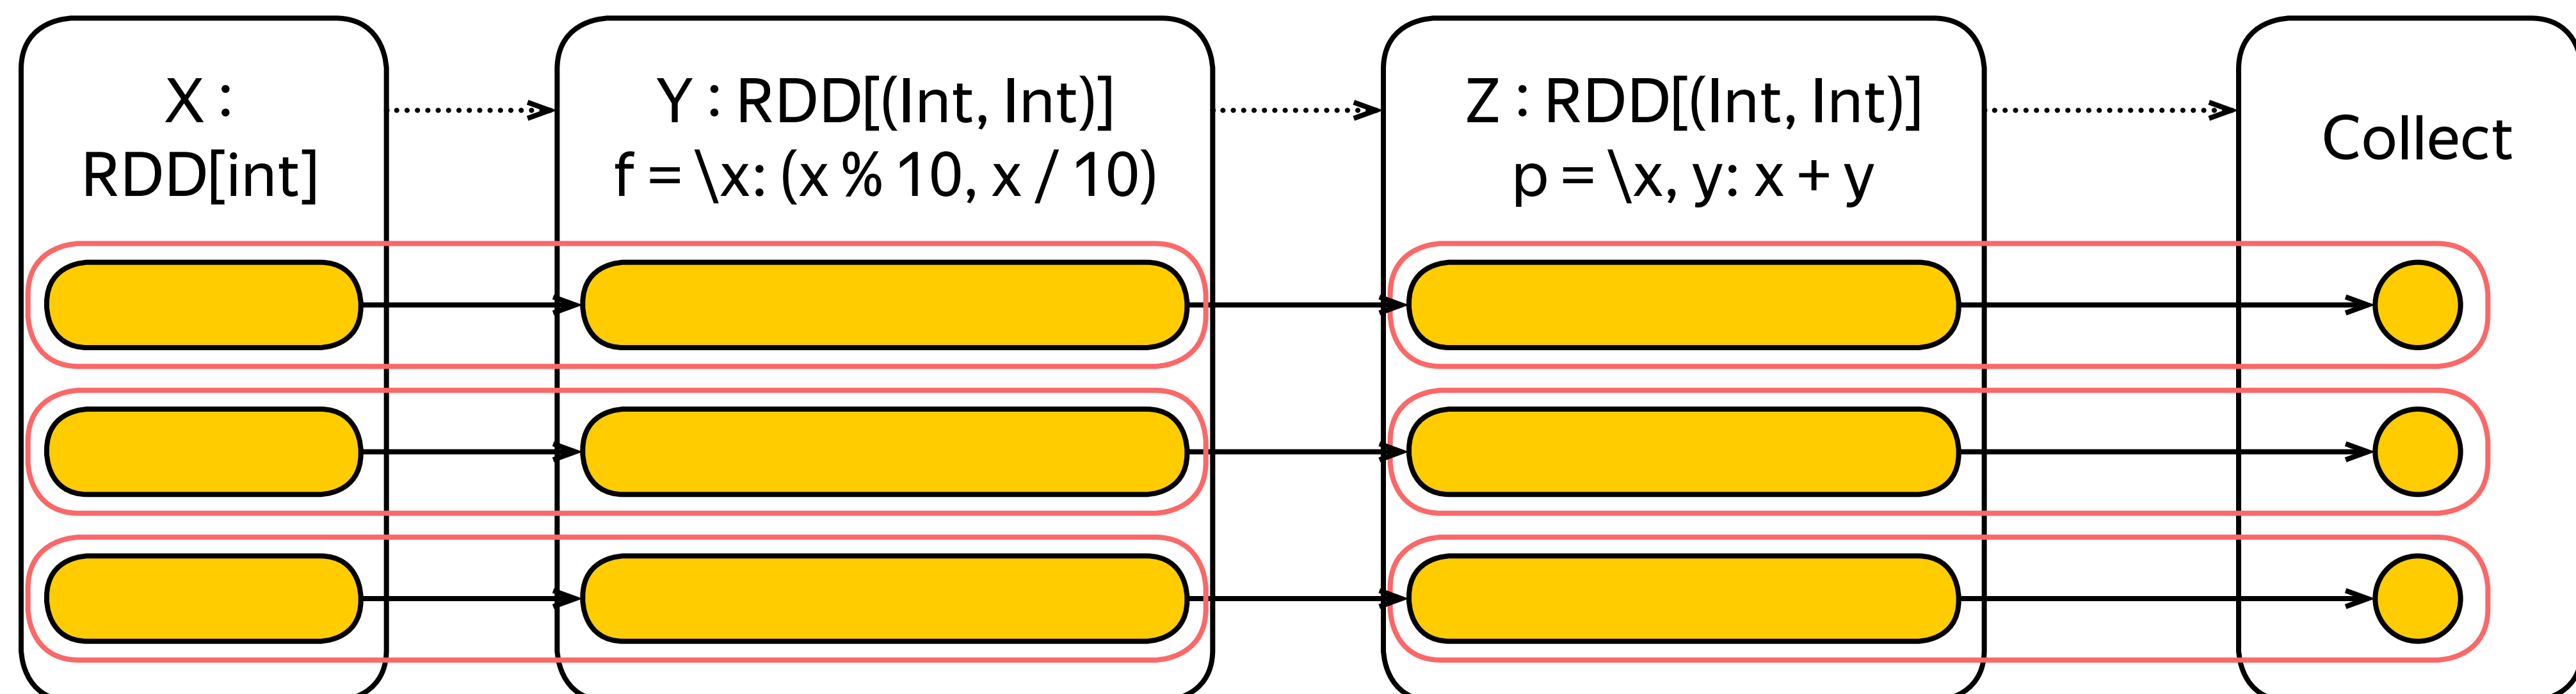


# Jobs, stages, tasks (example)

›  $Z = X$

```
.map(lambda x: (x % 10, x / 10))  
.reduceByKey(lambda x, y: x + y)  
.collect()
```

1. Invoking an action...
2. ...spawns the job...
3. ...that gets divided into the stages by the job scheduler...
4. ...and tasks are created for every job stage.



# Jobs, stages, tasks

- › **Job stage** is a pipelined computation spanning between materialization boundaries
  - › not immediately executable
- › **Task** is a job stage bound to particular partitions
  - › immediately executable

# Jobs, stages, tasks

- › **Job stage** is a pipelined computation spanning between materialization boundaries
- › **Task** is a job stage bound to particular partitions
- › Materialization happens when reading, shuffling or passing data to an action
  - › narrow dependencies allow pipelining
  - › wide dependencies forbid it

# SparkContext – other functions

- › Tracks liveness of the executors
  - › required to provide fault-tolerance
- › Schedules multiple concurrent jobs
  - › to control the resource allocation within the application
- › Performs dynamic resource allocation
  - › to control the resource allocation between different applications

# Summary

- › The SparkContext is the core of your application
- › The driver communicates directly with the executors
- › Execution goes as follows:  
Action → Job → Job Stages → Tasks
- › Transformations with narrow dependencies allow pipelining



**BigDATA**team