

Yandex

Caching & persistence

Quick reminder

- › RDDs are partitioned
- › Execution is build around the partitions
- › **Block** is a unit of input and output in Spark

Motivating example

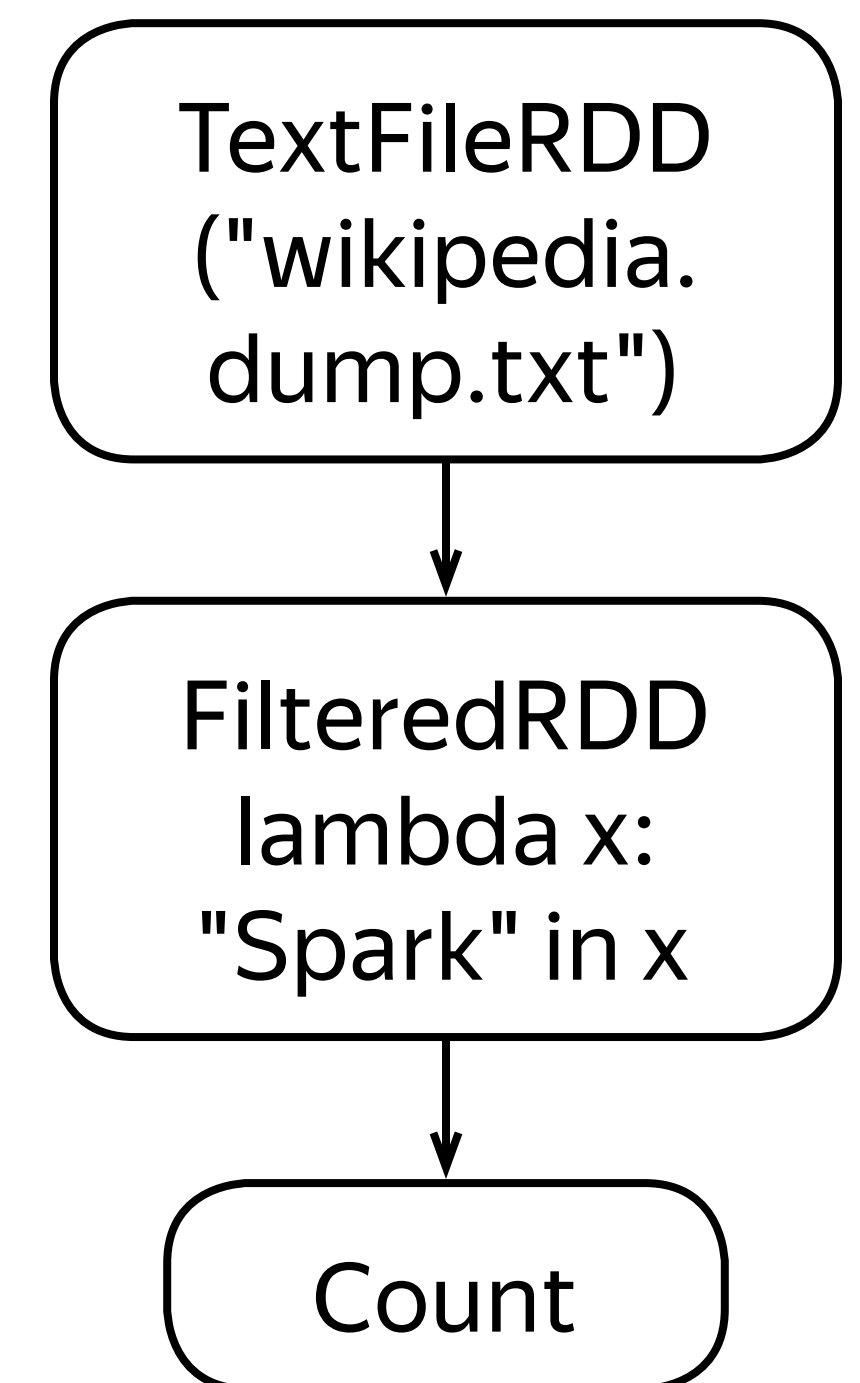
```
sc = SparkContext(...)
```

Motivating example

```
sc = SparkContext(...)
wiki = sc.textFile("wikipedia.dump.txt")

spark_articles = wiki.filter(
    lambda x: "Spark" in x)

print(spark_articles.count())
```

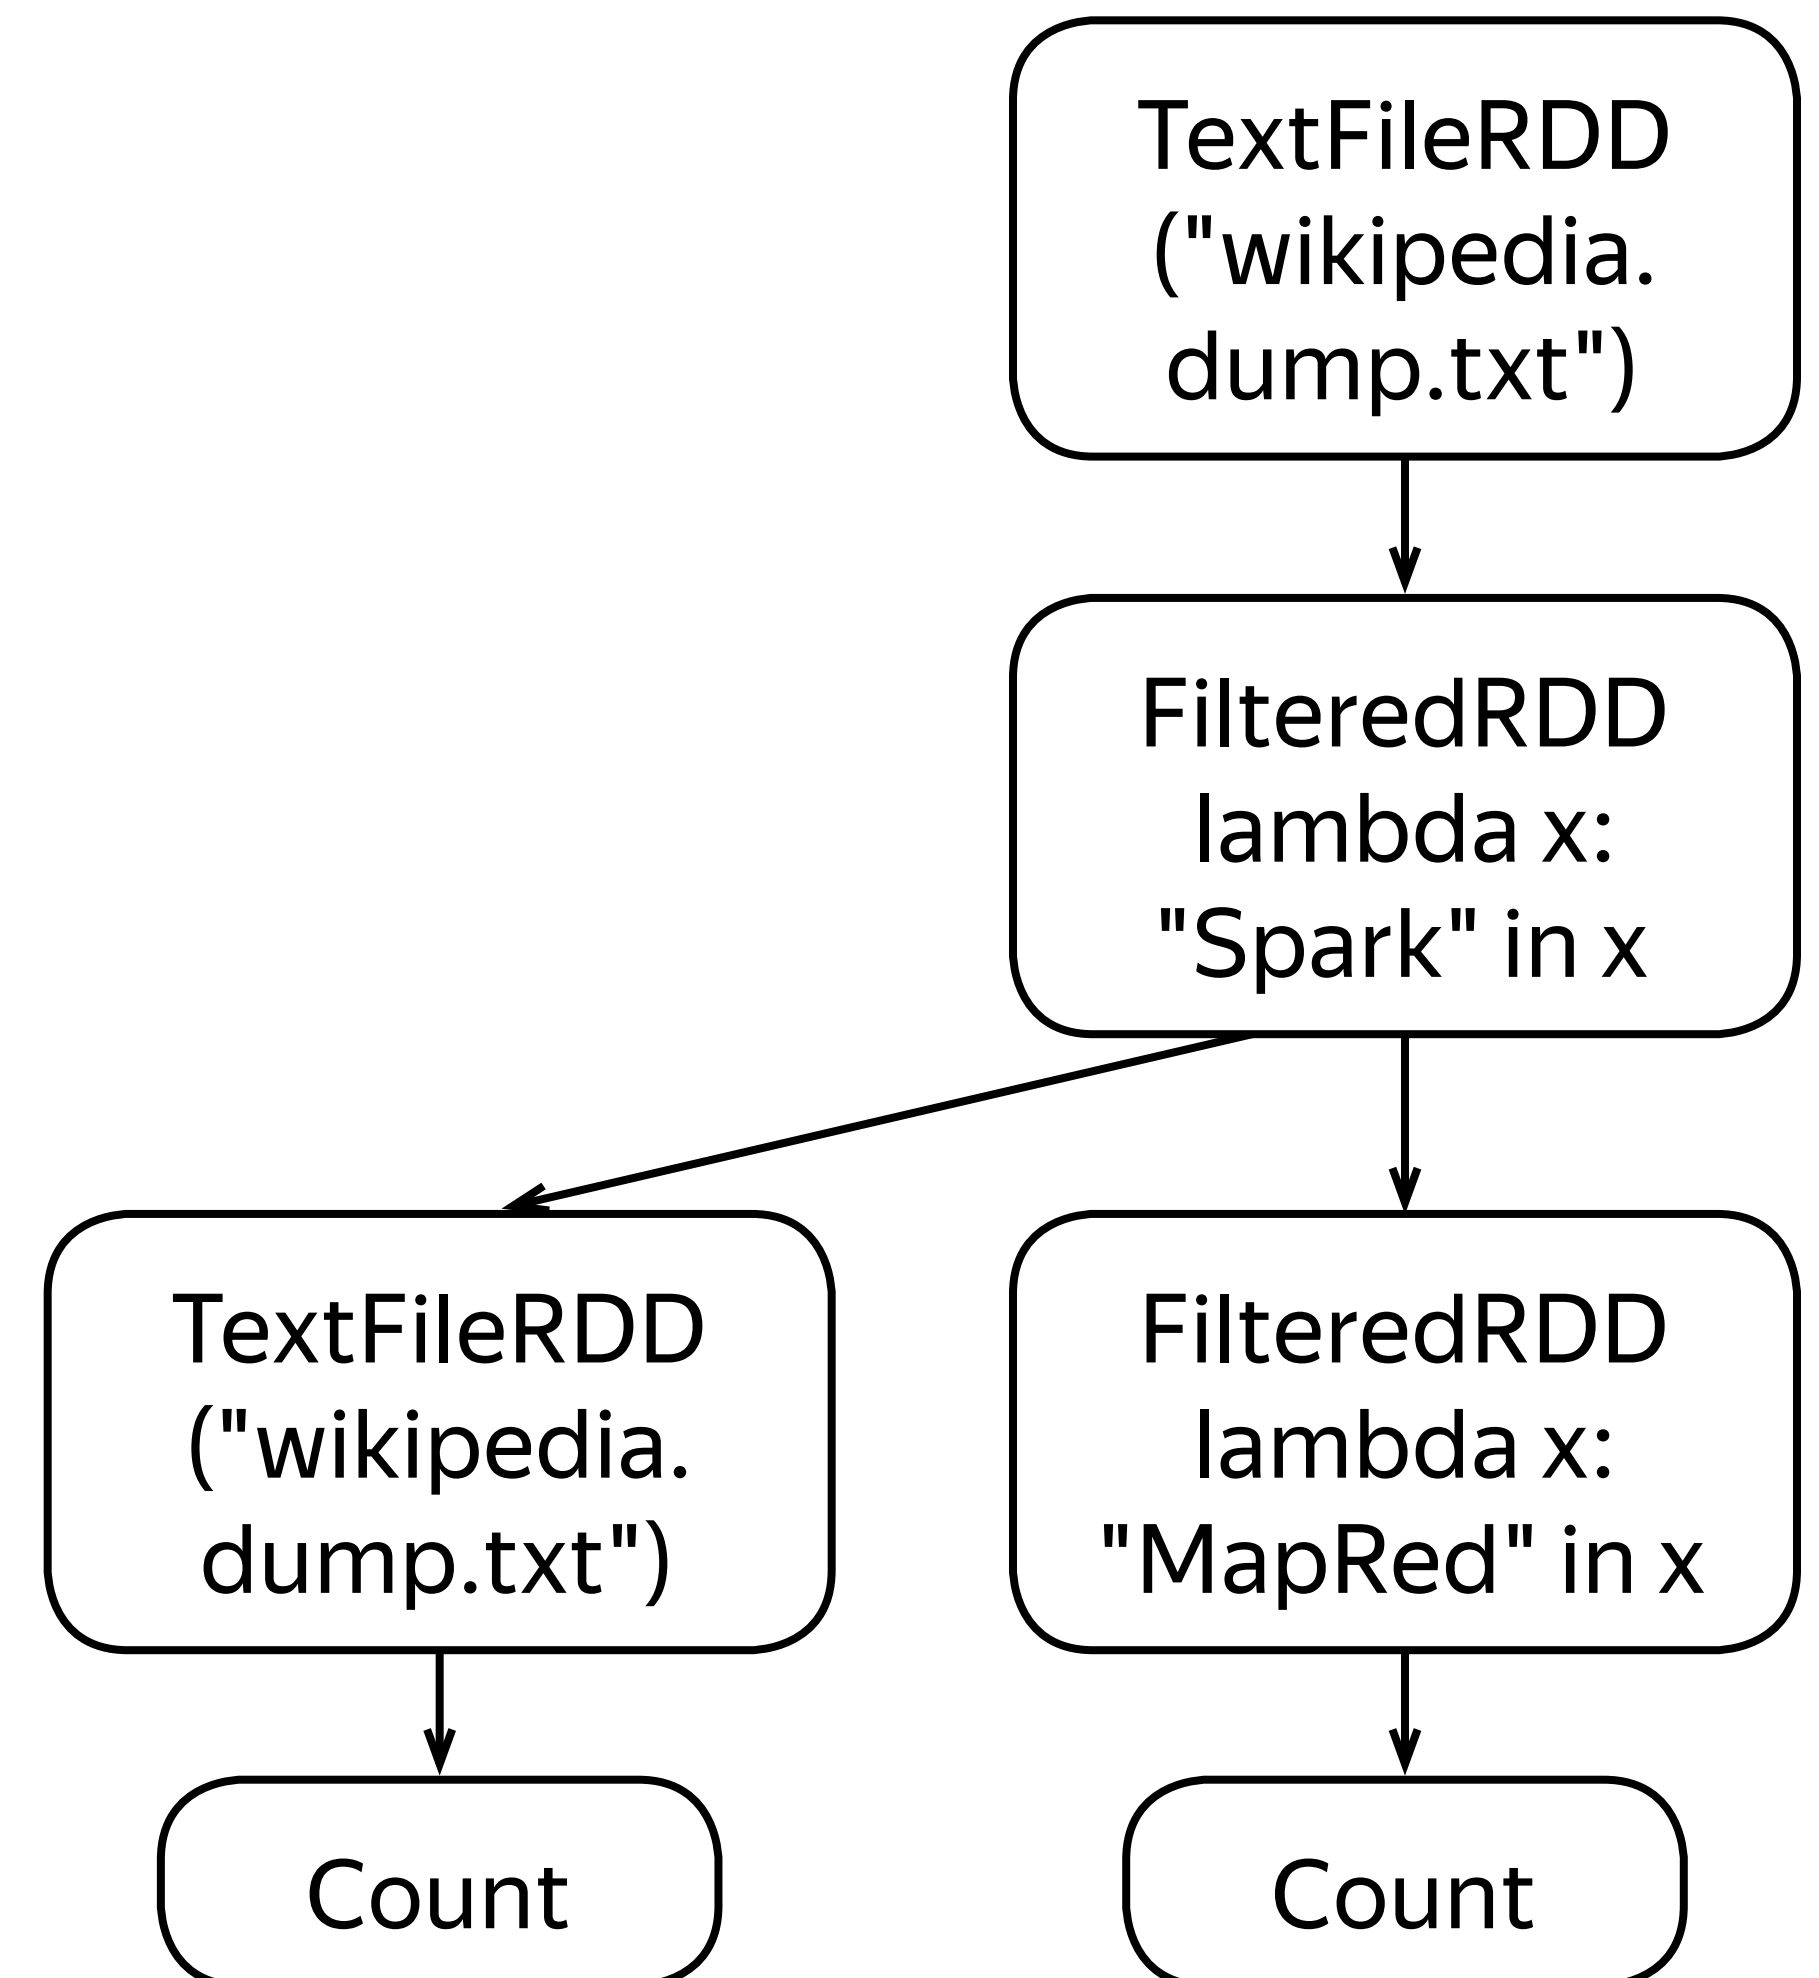


Motivating example

```
sc = SparkContext(...)
wiki = sc.textFile("wikipedia.dump.txt")

spark_articles = wiki.filter(
    lambda x: "Spark" in x)

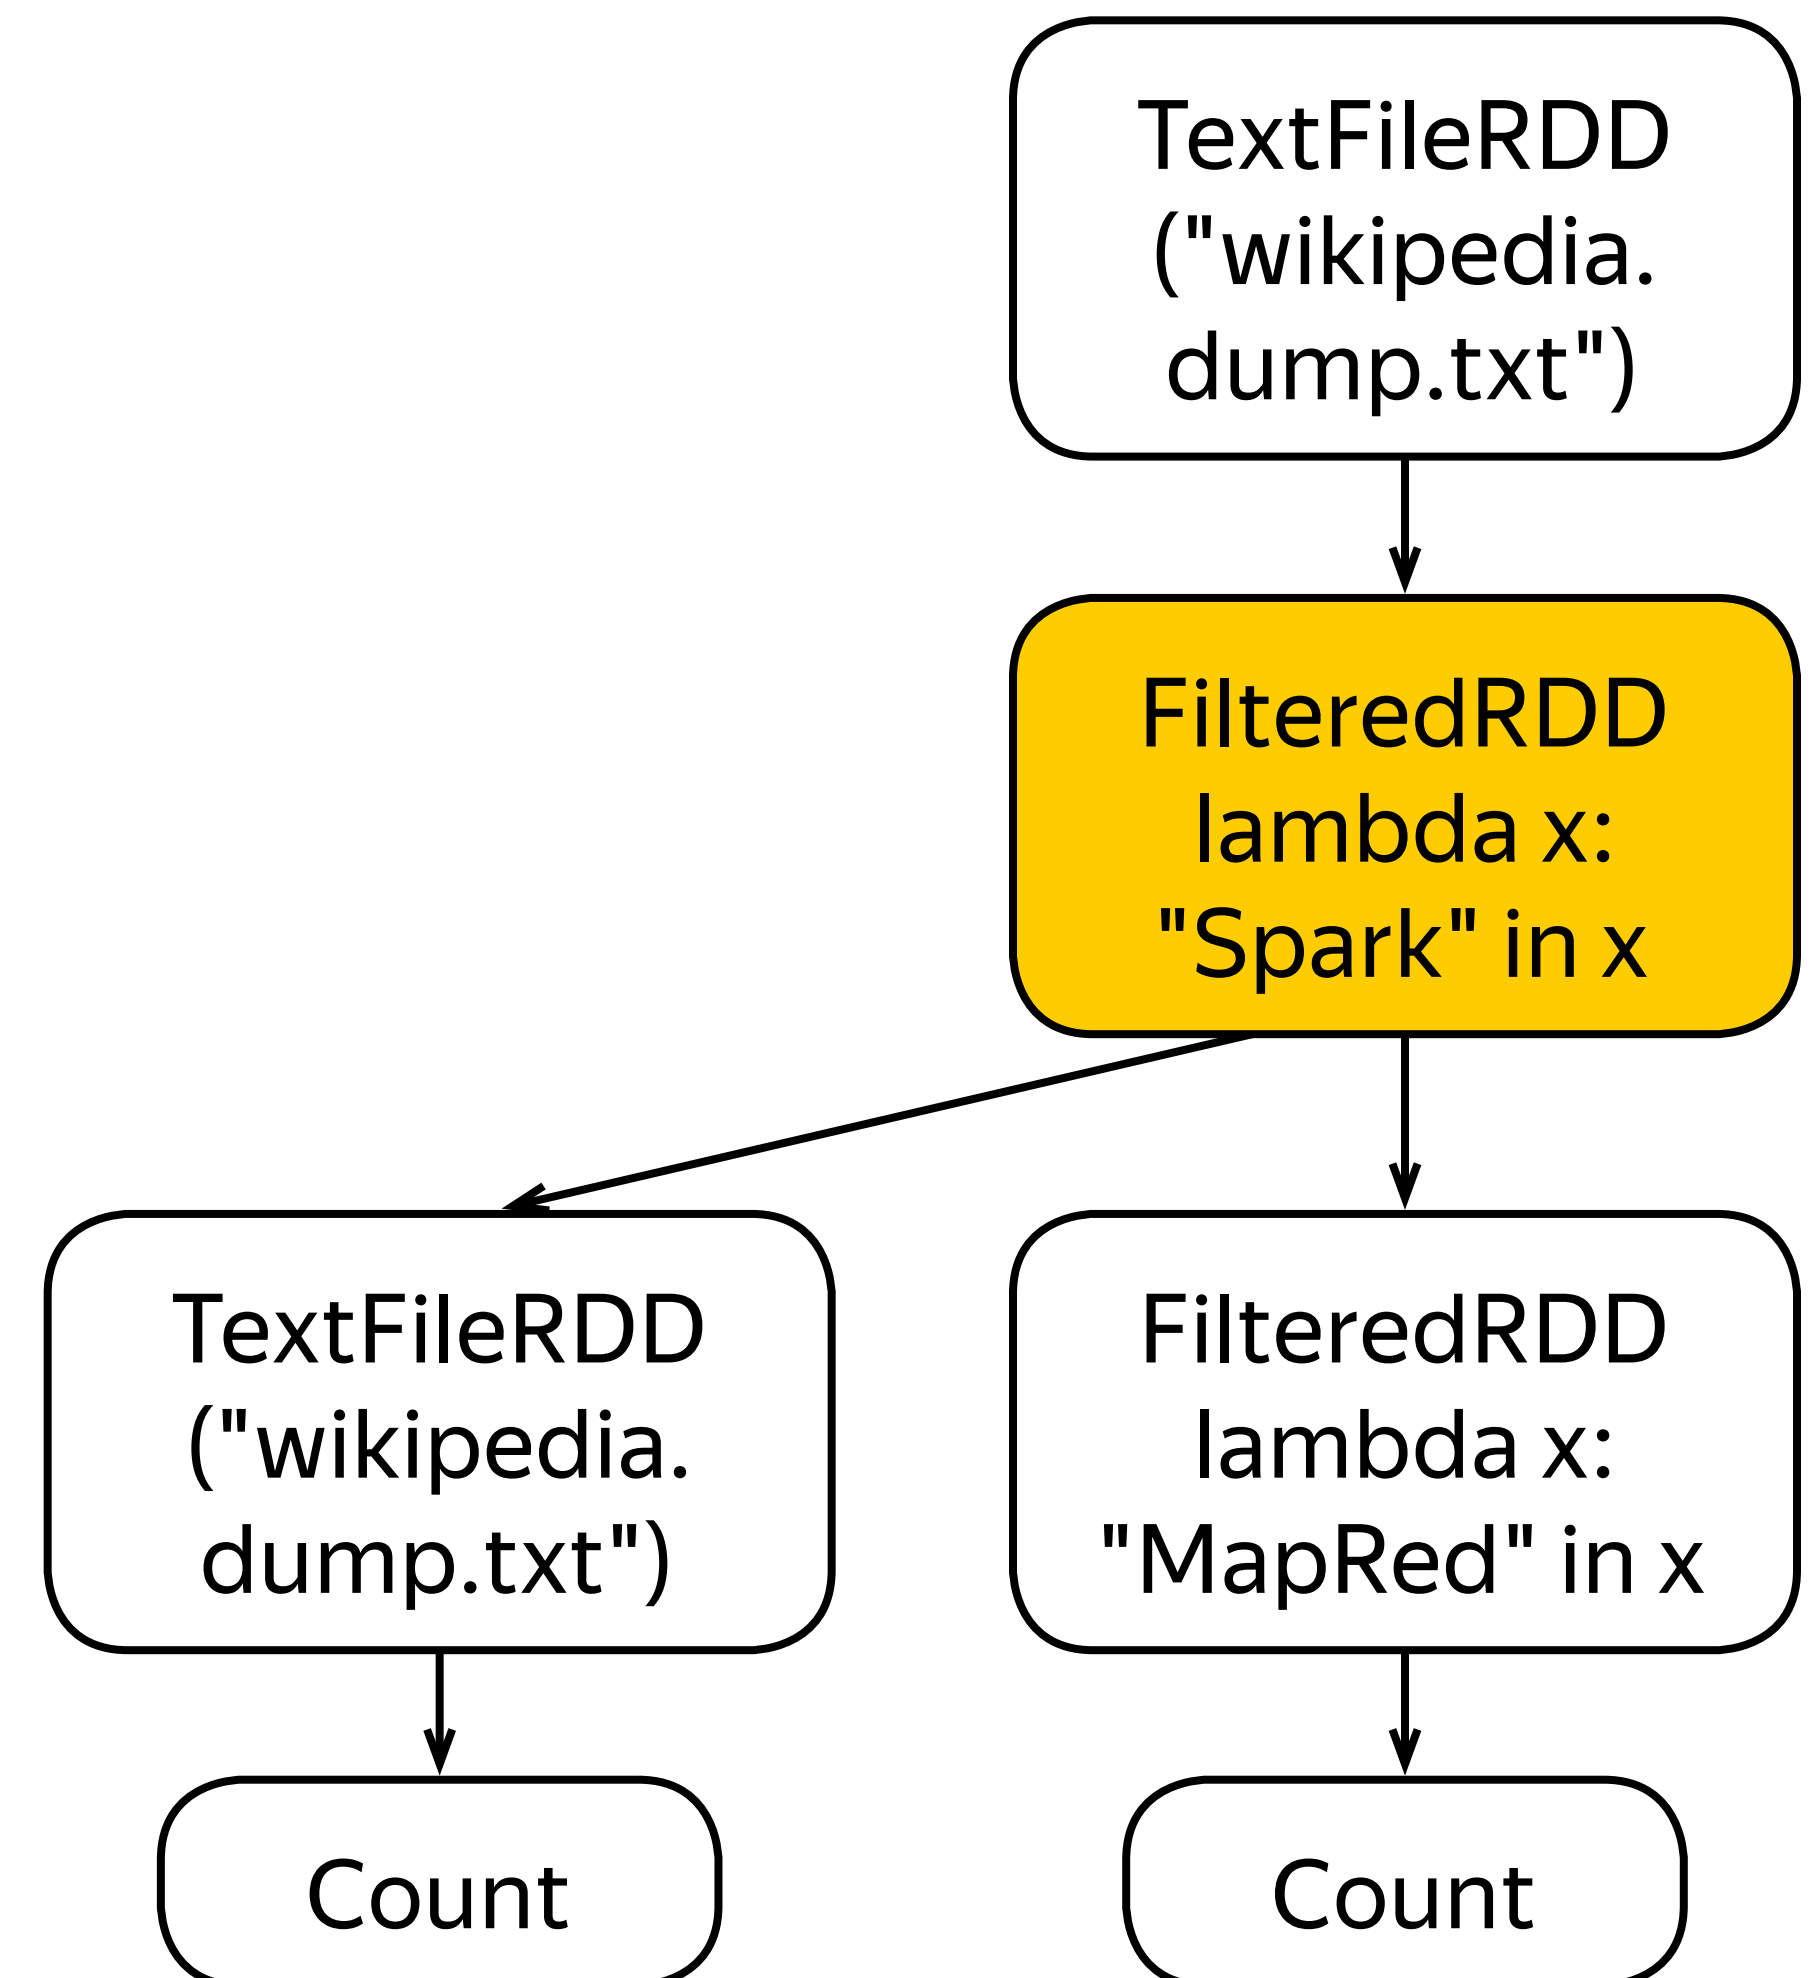
hadoop_articles = spark_articles.filter(
    lambda x: "Hadoop" in x)
mapreduce_articles = spark_articles.filter(
    lambda x: "MapRed" in x)
print(hadoop_articles.count())
print(mapreduce_articles.count())
```



Motivating example

```
sc = SparkContext(...)
wiki = sc.textFile("wikipedia.dump.txt")

spark_articles = wiki.filter(
    lambda x: "Spark" in x)
spark_articles.cache()
hadoop_articles = spark_articles.filter(
    lambda x: "Hadoop" in x)
mapreduce_articles = spark_articles.filter(
    lambda x: "MapRed" in x)
print(hadoop_articles.count())
print(mapreduce_articles.count())
```



Controlling persistence level

› `rdd.persist(storageLevel)`

› sets RDD's storage to persist across operations after it is computed for the first time

› `storageLevel` is a set of flags controlling the persistence, typical values are

`DISK_ONLY`

– save the data to the disk,

`MEMORY_ONLY`

– keep the data in the memory

`MEMORY_AND_DISK`

– keep the data in the memory; when out of memory – save it to the disk

`DISK_ONLY_2`, `MEMORY_ONLY_2`, `MEMORY_AND_DISK_2`

– same as about, but make two replicas

› `rdd.cache() = rdd.persist(MEMORY_ONLY)`

Controlling persistence level

› `rdd.persist(storageLevel)`

› sets RDD's storage to persist across operations after it is computed for the first time

› `storageLevel` is a set of flags controlling the persistence, typical values are

`DISK_ONLY`

– save the data to the disk,

`MEMORY_ONLY`

– keep the data in the memory

`MEMORY_AND_DISK`

– keep the data in the memory; when out of memory – save it to the disk

`DISK_ONLY_2`, `MEMORY_ONLY_2`, `MEMORY_AND_DISK_2`

– same as about, but make two replicas ← improves failure recovery times!

› `rdd.cache() = rdd.persist(MEMORY_ONLY)`

Best practices

Best practices

- › For interactive sessions
 - › cache preprocessed data

Best practices

- › For interactive sessions
 - › cache preprocessed data
- › For batch computations
 - › cache dictionaries
 - › cache other datasets that are accessed multiple times

Best practices

- › For interactive sessions
 - › cache preprocessed data
- › For batch computations
 - › cache dictionaries
 - › cache other datasets that are accessed multiple times
- › For iterative computations
 - › cache static data
- › And do benchmarks!

Summary

- › Performance may be improved by persisting data across operations
 - › in interactive sessions, iterative computations and hot datasets
- › You can control the persistence of a dataset
 - › whether to store in the memory or on the disk
 - › how many replicas to create

BigDATAteam